

Report on project 1

Ekaterina Ilin and Isabelle Gauger

September 7, 2014

Contents

1	The algorithm	2
2	Implementation and results	4
3	The relative error	5
4	Comparison of execution time with the LU-decomposition	5

1 The algorithm

In order to solve the one-dimensional Poisson equation

$$-u''(x) = f(x) \quad (1)$$

with Dirichlet boundary conditions in the interval $(0, 1)$ we rewrite the latter as a set of linear equations by discretizing the problem. In this way we obtain a set of n grid points with the gridwidth $h = 1/(n + 1)$. Then we approximate the second derivative $u''(x)$ with

$$-\frac{-v_{i+1} - v_{i-1} + 2v_i}{h^2} = f_i \quad \text{for } i = 1, \dots, n \quad (2)$$

If we now multiply these equations with h^2 and consider the boundary conditions $v_0 = v_{n+1} = 0$ we obtain a set of equations that can be written as follows:

$$\begin{array}{cccccccl} 2v_1 & -v_2 & & & & & = h^2 f_1 \\ -v_1 & +2v_2 & -v_3 & & & & = h^2 f_2 \\ & -v_2 & +2v_3 & -v_4 & & & = h^2 f_3 \\ & \dots & & & & & \\ & & & -v_{n-1} & +2v_n & & = h^2 f_n \end{array} \quad (3)$$

From these n equations we can easily derive the following $n \times n$ -matrix equation:

$$\begin{pmatrix} 2 & -1 & 0 & \dots & \dots & 0 \\ -1 & 2 & -1 & 0 & \dots & \dots \\ 0 & -1 & 2 & -1 & 0 & \dots \\ & \dots & \dots & \dots & \dots & \dots \\ 0 & \dots & & -1 & 2 & -1 \\ 0 & \dots & & 0 & -1 & 2 \end{pmatrix} \begin{pmatrix} v_1 \\ v_2 \\ \dots \\ \dots \\ \dots \\ v_n \end{pmatrix} = \begin{pmatrix} h^2 f_1 \\ h^2 f_2 \\ \dots \\ \dots \\ \dots \\ h^2 f_n \end{pmatrix}. \quad (4)$$

A more general form of the above is:

$$\begin{pmatrix} a_0 & b_0 & 0 & \dots & \dots & \dots \\ c_1 & a_1 & b_1 & \dots & \dots & \dots \\ & c_2 & a_2 & b_2 & \dots & \dots \\ & \dots & \dots & \dots & \dots & \dots \\ & & & c_{n-2} & a_{n-2} & b_{n-2} \\ & & & c_{n-1} & a_{n-1} & \end{pmatrix} \begin{pmatrix} v_0 \\ v_1 \\ \dots \\ \dots \\ \dots \\ v_{n-1} \end{pmatrix} = \begin{pmatrix} w_0 \\ w_1 \\ \dots \\ \dots \\ \dots \\ w_{n-1} \end{pmatrix}. \quad (5)$$

Since the Gaussian elimination would of course lead to the correct results here, the execution time can be easily reduced from $\sim n^3$ to $\sim n$ by applying an algorithm that no longer requires the matrix but uses the three diagonals as arrays. In other words we consider that the rest of the matrix is 0 everywhere except for these diagonals which the brute force Gaussian elimination way does not take in account. The following steps have then to be taken:

1. The three diagonals are stored in arrays $a[]$, $b[]$, and $c[]$, as well as the right side of the equation is stored in an array $w[]$ of the size n . $c[0]$ and $b[n - 1]$ are set to 0.
2. Then the entries in $a[]$ are substituted recursively by

$$\tilde{a}[0] = a[0], \quad \tilde{a}[i] = a[i] - b[i - 1] \frac{c[i]}{\tilde{a}[i - 1]} \quad (6)$$

This requires $3 \cdot (n - 1)$ floating point operations for we obtain a division, a subtraction and a multiplication for each substitution.

3. Accordingly $w[]$ is substituted by

$$\tilde{w}[0] = w[0], \quad \tilde{w}[i] = w[i] - \tilde{w}[i - 1] \frac{c[i]}{\tilde{a}[i - 1]} \quad (7)$$

This only requires $2 \cdot (n - 1)$ flops for we already did the division $\frac{c[i]}{\tilde{a}[i - 1]}$ during the substitution above.

4. Finally backward substitution is used to gain the result for the unknown vector v which is stored in another array $v[]$:

$$v[n - 1] = \frac{\tilde{w}[n - 1]}{\tilde{a}[n - 1]}, \quad v[i] = \frac{\tilde{w}[i] - b[i] \cdot v[i + 1]}{\tilde{a}[i - 1]} \quad (8)$$

This operation results in another $3 \cdot (n - 1) + 1$ flops for we have again a subtraction, a multiplication and a division for each resubstitution plus a division for the first element.

In sum the algorithm needs $8 \cdot (n - 1) + 1$ floating point operations to solve the general matrix equation 5. If we now go back to the specific matrix 4 we can simplify the algorithm once more by tuning the above one to our needs. For this we go through the previous algorithm and insert the given values for the entries of the vectors $a[]$, $b[]$ and $c[]$:

1. The entries in the diagonals now are:

$$a[i] = 2 \quad \text{for } i = 0, \dots, n-1, \quad b[i] = -1 \quad \text{for } i = 0, \dots, n-2, \quad c[i] = -1 \quad \text{for } i = 1, \dots, n-1 \quad (9)$$

In this way the array computed in 6 becomes a static vector that can be computed recursively once and then stored somewhere for it does not depend on the right side of the equation:

$$\tilde{a}[0] = a[0], \quad \tilde{a}[i] = a[i] - b[i - 1] \frac{c[i]}{\tilde{a}[i - 1]} = a[i] - \frac{1}{\tilde{a}[i - 1]} \quad (10)$$

Therefore this does require $2 \cdot (n - 1)$ flops - but only once, so we can compute $\tilde{a}[]$ once for a very large n and then make use of it for any given $w[]$.

2. Latter then is to be substituted as follows

$$\tilde{w}[0] = w[0], \quad \tilde{w}[i] = w[i] - \tilde{w}[i - 1] \frac{c[i]}{\tilde{a}[i]} = w[i] + \frac{\tilde{w}[i - 1]}{\tilde{a}[i]} \quad (11)$$

This requires $2 \cdot (n - 1)$ floating point operations.

3. The backward substitution can then be simplified, too:

$$v[n - 1] = \frac{\tilde{w}[n - 1]}{\tilde{a}[n - 1]}, \quad v[i] = \frac{\tilde{w}[i] - b[i] \cdot v[i + 1]}{\tilde{a}[i]} = \frac{\tilde{w}[i] + v[i + 1]}{\tilde{a}[i]} \quad (12)$$

This operation leads to another $2 \cdot (n - 1) + 1$ flops.

In sum we obtain $4 \cdot (n - 1) + 1$ floating point operations which is about twice as fast as the general algorithm for tridiagonal matrices. If we used standard Gaussian elimination here we would have ended up with an execution time $O(\frac{2}{3}n^3)$. An application of the LU-decomposition could have been applied but would not have been worth the effort since - in this case - there is only one vector w to be solved for. The standard LU-decomposition itself requires $O(\frac{2}{3}n^3)$ floating point operations. For solving the set of equations then we would have needed another $O(n^2)$ operations which in sum is even worse than the standard Gaussian elimination.

2 Implementation and results

Possible implementations of both algorithms for tridiagonal matrices are shown below, as well as the plotted results of the computations with $n = 10$, 100 and 1000 (see figures 1, 2 and 3) respectively.

solvers for tridiagonal matrices

```
double* tridiagonal (double*a, double*b, double*c, double*w, int n)
//solves the general tridiagonal matrix
{
    double* v;
    v = new double[n];

    for (int i=1; i<n; i++)
    {
        double temp = c[i]/a[i-1];           //saves n-1 flops
        a[i] -= b[i-1]*temp;                 //1.
        w[i] -= w[i-1]*temp;                 //2.
    }

    v[n-1] = w[n-1]/a[n-1];                 //3.

    for (int i=n-2; i>=0; i--)
    {
        v[i] = (w[i]-b[i]*v[i+1])/a[i];
    }
    return v;
}

double* tridiagonaldiff (double* a, double* w, int n)
//solves the discretized Poisson equation, here *a is a pointer to the
//static array (see 1.)
{
    double* v;
    v = new double[n];

    for (int i=1; i<n; i++)
    {
        w[i] += w[i-1]/a[i-1];               //2.
    }

    v[n-1] = w[n-1]/a[n-1];                 //3.
```

```

    for (int i = n-2; i >= 0; i--)
    {
        v[i] = (w[i]+v[i+1])/a[i];
    }
    return v;
}

```

3 The relative error

The relative error ϵ_i for the data set $i = 1, \dots, n$ is given by:

$$\epsilon_i = \log_{10} \left(\left| \frac{v_i - u_i}{u_i} \right| \right) \quad (13)$$

We started with $n = 10$ steps and then increased the number of steps by factor 10 until $n = 10^6$. The maximum values of the relative error for the different step lengths h are shown in table 1 and plotted in figure 4.

When we increase the number of steps what means to reduce the step length the relative error decreases linearly (with the \log_{10} scale) until 10^5 steps. When we increase the number of steps further to $n = 10^6$ the relative error becomes larger again. That shows that there is a limit to how small we can make the step length, before we run into problems with loss of precision. This also means that we can only reduce the relative error of our computed results to a certain limit.

4 Comparison of execution time with the LU-decomposition

Using the LU-decomposition functions **ludcmp** and **lubskb** given in lib.cpp to solve the Poisson equation we obtain results that are nearly identical to the results from our own algorithm. Table 2 shows the relative errors between our solutions and the results from the LU decomposition. Compared to the relative error we obtain from the comparison to the reference solution one can see that the quality of both results is approximately the same since the relative error between our solution and the LU-solution is at least smaller by 3 magnitudes than the relative error we obtain between our solution and the reference.

The comparison of execution time between the two algorithms vividly shows the importance of tuning algorithms to specific needs in order to get rid of unnecessary execution time. Table

Table 1: max value of the relative error for different step lengths

$\log_{10}(h)$	ϵ_i	n
-1.041392685	-1.179697782	10
-2.004321374	-3.088036832	100
-3.000434077	-5.080051538	1000
-4.000043427	-7.079270511	10000
-5.000004343	-8.847801518	100000
-6.000000434	-8.05486036	1000000

Table 2: maximum relative errors between our solution and the LU-decomposition

n	$\epsilon_{max,n}$
10^1	-15.27
10^2	-13.91
10^3	-12.41

Table 3: comparison of execution times between our solution and the LU-decomposition

n	t in s for tridiagonaldiff	t in s for ludcmp + lubksb
10^1	0	0
10^2	0	0.002
10^3	0	2.393
10^4	0	estimated 40min, too long to perform
10^5	0.003	estimated 11 days
10^6	0.033	-
10^7	0.359	-

3 illustrates the growth of the execution time with increasing n . The resolution of the timer used is about 1ms. From this we can conclude that the execution time for our linear algorithm is below 1ms for step numbers up to $n = 10^4$. For $n = 10^5$ we obtain an execution time of a few milliseconds and for $n = 10^6$ the time to perform the algorithm was approximately ten times longer which is exactly what we discussed in the previous paragraphs.

The LU-decomposition algorithm also behaves as expected. For $n = 10$ the execution time is shorter than the precision of the clock. But for $n = 10^2$ we already get an execution time of 2ms and for $n = 10^3$ the time rises by the factor 10^3 . This verifies the previously discussed number of flops needed for the performance of this algorithm which is $O(n^3)$. According to this the execution times for the LU-decomposition for $n = 10^4$ and $n = 10^5$ should rise to about 40min and 11 days respectively. This is far beyond acceptable.




Figure 1: plot with $n = 10$




Figure 2: plot with $n = 100$




Figure 3: plot with $n = 1000$




Figure 4: max value of the relative error for different step lengths